

PROJECT REPORT DOCUMENTATION

Date: 15/09/2026

Submitted By:

Name: Satya Nachika
Department: MSc. Physics [IITM]

Submitted To: GENZ EDUCATEWING Team

(1). Project Name:

Sentiment Analysis Using BERT (Bidirectional Encoder
Representations From Transformers)

1. Abstract

Sentiment analysis is a Natural Language Processing (NLP) technique used to identify and classify the emotional tone expressed in textual data. With the rapid growth of digital platforms such as social media, online shopping websites, and customer feedback systems, large amounts of text data are generated every day. Automatically analyzing this information helps organizations understand user opinions and make better decisions.

In this project, a deep learning-based sentiment analysis system is developed using the BERT (Bidirectional Encoder Representations from Transformers) model. The main objective is to classify text reviews into three sentiment categories: **Positive, Negative, and Neutral**.

The dataset is first processed through different preprocessing techniques such as removal of unwanted characters, URLs, punctuation, stop words, and lemmatization. The cleaned text is converted into numerical representations using the BERT tokenizer. A pre-trained BERT model (Bert-base-uncased) is then fine-tuned using the prepared dataset with the Hugging Face Transformers framework.

The performance of the model is evaluated using evaluation metrics such as accuracy, precision, recall, F1-score, and confusion matrix. The results show that BERT can effectively capture contextual information from text and provide accurate sentiment classification.

2. Introduction

The increasing availability of digital text data has created a need for automatic methods to analyze human opinions and emotions. Sentiment analysis provides a computational approach to determine whether a given text expresses a positive, negative, or neutral opinion.

Traditional machine learning approaches such as Naive Bayes, Support Vector Machines (SVM), and Logistic Regression require manual feature extraction techniques like Bag-of-Words and TF-IDF.

Recent developments in deep learning introduced transformer-based models that can understand the relationship between words in a sentence. BERT, developed by Google Research, is one of the most successful transformer architectures for NLP applications. Unlike traditional models, BERT uses bidirectional attention to understand words based on their surrounding context.

In this project, BERT is fine-tuned for sentiment classification to develop an efficient system capable of analyzing textual reviews.

3. Objectives

The main objectives of this project are:

1. To develop a sentiment analysis model using the BERT transformer architecture.
 2. To perform text preprocessing and data cleaning.
 3. To convert textual information into numerical form using BERT tokenizer.
 4. To fine-tune a pre-trained BERT model for sentiment classification.
 5. To classify reviews into Positive, Negative, and Neutral categories.
 6. To evaluate the model performance using standard evaluation metrics.
 7. To analyze model performance through visualization techniques.
-

4. Dataset Description

The dataset contains text reviews and their corresponding sentiment labels.

The main features of the dataset are:

Feature	Description
Text	Original user review
Sentiment	Sentiment category
Clean-text	Preprocessed review text
Label	Numerical representation of sentiment

The sentiment classes considered in this project are:

- Negative
- Neutral
- Positive

The dataset is divided into training and testing sets using an 80:20 ratio. The training dataset is used for fine-tuning the BERT model, while the testing dataset is used for performance evaluation.

5. Methodology

The complete workflow of the project consists of the following steps:

5.1 Data Preprocessing

Raw text data usually contains unnecessary information such as URLs, special characters, and common words that do not contribute significantly to sentiment classification. Therefore, preprocessing is performed before model training.

The preprocessing steps include:

Lowercase Conversion

All text samples are converted into lowercase to maintain consistency.

Example:

Before:

The Product is GREAT

After:

the product is great

Removal of URLs and Special Characters

Unwanted links, symbols, and punctuation marks are removed from the text.

Stop Word Removal

Common English words such as “the”, “is”, and “a” are removed.

Lemmatization

Words are converted into their base form.

Example:

playing → play


6. BERT Model Implementation

The model used in this project is:

BERT-base-uncased

BERT consists of transformer encoder layers that learn contextual relationships between words.

The processing flow is:



The tokenizer converts text into tokens and adds special tokens such as:

[CLS] sentence [SEP]

The transformer encoder learns the meaning of words based on their context. The final classification layer predicts the sentiment category.

7. Model Training

The pre-trained BERT model is fine-tuned using the prepared sentiment dataset.

Training parameters:

Parameter	Value
Model	Bert-base-uncased
Learning Rate	2×10^{-5}
Epochs	3
Batch Size	16
Maximum Sequence Length	128
Optimizer	AdamW

During training, the model learnt the relationship between input text and sentiment labels.

8. Model Evaluation

The performance of the trained model is measured using different evaluation metrics.

Accuracy: Accuracy represents the percentage of correctly classified samples.

Accuracy = 0.3333333333333333

Precision: Precision measures how many predicted positive results are actually correct.

Precision = 0.1111111111111111

Recall: Recall represents the ability of the model to identify all relevant samples.

Recall = 0.3333333333333333

F1 Score: F1-score provides a balance between precision and recall.

[F1 = 0.16666666666666666]

Confusion Matrix: The confusion matrix provides class-wise information about correct and incorrect predictions.

9. Results and Analysis

The trained BERT model is evaluated using the test dataset. The following analyses are performed:

Performance Metrics

The model performance is measured using:


- Accuracy
- Precision
- Recall
- F1-score

Visualization

Different graphs are generated:


1. Confusion Matrix

Shows correct and incorrect predictions for each sentiment class.




2. Performance Comparison Graph

- Compares accuracy, precision, recall, and F1-score.



3. Training and Validation Loss Curves


- Shows the learning character of the model during training.



The decreasing loss values indicate that the model successfully learns patterns from the training data

4. Other graphs to classify reviews into Positive and Negative.


(a): This is the Positive Reviews Word Cloud figure




(b): This is the Negative Reviews Word Cloud figure




(c): This is The Review length Boxplot Figure.




(d) This is the Review length Distribution Figure.



(e) This is the Sentiment Distribution Figure.



(Pie-Chart)



(Histogram)

10. Applications

The developed sentiment analysis system has several applications as follows given below:

Product Review Analysis: Companies can automatically analyze customer feedback.

Social Media Monitoring: Organizations can understand public opinions from social media platforms.

Customer Support: Feedback can be automatically classified and prioritized.

Market Research: Companies can study the consumer's like, dislike of the product and preferences.

Brand Reputation Management: Businesses can monitor online opinions about their products.

11. Advantages

The major advantages of the BERT-based approach are:

- Better understanding of word context.
 - High classification accuracy.
 - Requires less manual feature extraction.
 - Can be adapted to various NLP tasks.
 - Uses transfer learning from a pre-trained model.
-

12. Limitations

Although BERT provides excellent performance, it has some limitations:

- Requires high computational resources.
 - Training time is longer compared to traditional machine learning models.
 - Large model size requires more memory.
 - Performance depends on dataset quality.
-

13. Future Scope

The project can be improved in several ways:

1. Using advanced transformer models such as RoBERTa and DeBERTa.
 2. Developing a real-time sentiment analysis application.
 3. Extending the system for multilingual sentiment classification.
 4. Deploying the model using cloud platforms.
 5. Using explainable AI methods to understand model predictions.
-

14. Conclusion

In this project, a sentiment analysis system was successfully developed using the BERT transformer model. The complete pipeline including data preprocessing, BERT tokenization, model fine-tuning, evaluation, and prediction was implemented.

The BERT model effectively learned contextual information from text data and classified reviews into positive, negative, and neutral categories. The evaluation using accuracy, precision, recall, F1-score, and confusion matrix demonstrated the effectiveness of transformer-based models for sentiment analysis.

This project shows the importance of deep learning and transformer architectures in solving modern Natural Language Processing problems. The developed model can be further extended for real-world applications such as customer feedback analysis, social media monitoring, and automated opinion mining.
